

# Math 578 Final

Luke Wukmer

Fall 2016

## 1 The Problem

We use Multigrid-Preconditioned Conjugate Gradient Method (henceforth referred to as **MGCG**) to solve a particular problem:

$$Ax = b$$

where

$$A \in \mathbb{R}^{2^M \times 2^M}, \quad A_{ij} = \begin{cases} 3 & \text{if } j = i, \\ -1 & \text{if } j = i \pm 1 \\ -1 & \text{if } j = N - i + 1, i \neq \frac{N}{2}, \frac{N}{2} + 1 \end{cases}, \quad (1)$$

and  $b$  is a  $N \times 1$  normalized one vector.

We will compare two methods of preconditioning via multigrid, defined here:

### 1.1 Scenario #1

In "scenario #1" (referred to henceforth as **SC1**), we precondition via multigrid using the following interpolation matrix (between levels  $l$  and  $l - 1$ ) (refer to equation #2 in the project spec):

```
1 def interpolation_matrix(M,L,el):
2     """
3     The interpolation matrix I_el between levels el and (el-1)
4     INPUT:
5     M -      refers to size of system 2^M x 2^M
6     L -      number of levels/grids
7     el-      interpolate between el and el-1
8
9     OUTPUT:
10    I_el     a 2^{ M + el - L } by 2^{M + el - (L + 1) }
11             interpolation matrix
12    """
13
14    n = 2**(M + el - (L+1))
15
16    # identity matrix but repeat each row twice
17    # (equivalent to given system)
18    return np.repeat(np.eye(n), 2, axis=0)
```

## 1.2 Scenario #2

Scenario #2, (SC2) is the same as the above, but we change precisely one level of interpolation: between the second and third finest levels. Refer to spec and the code below.

```
1 def interpolation_matrix_2(M,L,el):
2     """
3     The interpolation matrix I_el between levels el and (el-1)
4     for example if M = 2, this function would return
5         array([[ 1.,  0.,  0.,  0.],
6                [ 0.,  1.,  0.,  0.],
7                [ 0.,  0.,  1.,  0.],
8                [ 0.,  0.,  0.,  1.],
9                [ 0.,  0.,  0.,  1.],
10               [ 0.,  0.,  1.,  0.],
11               [ 0.,  1.,  0.,  0.],
12               [ 1.,  0.,  0.,  0.]])
13
14     for the given question, this is only to be used for the
15     (L-1) to (L-2)th level
16     """
17     if el != L-1:
18         return interpolation_matrix(M,L,el)
19     else:
20         n = 2*(M-2)
21         return np.concatenate((np.eye(n),np.flipud(np.eye(n))))
```

## 2 Implementation

### 2.1 The V-cycle

The following is a basic implementation of the multigrid "v-cycle."

```
1 def vcycle(l,b,e0, A, I, Lchol):
2
3     omega = 2/3
4     nu1 = 1
5     # base case
6     if l == 1:
7         e_base = linalg.solve(Lchol.T,linalg.solve(Lchol,b))
8         return e_base
9     else:
10         a = A[-(l-1)]
11         i = I[-(l-1)]
12         e = smooth(a, omega, nu1, b, e0)
13         # compute and restrict error
14         res = i.T @ (b - a@e)
15         # correct error
16         e = e + i @ vcycle(l-1,res, np.zeros_like(res), A,I,Lchol)
17         # smooth nu1 times on a x = b with initial guess e
18         e = smooth(a,omega,nu1,b,e)
19
20     return e
```

## 2.2 Smoothing function

Here is the smoothing function. As described in the docstring, this function was first tested for accuracy by applying it to a  $4 \times 4$  diagonally dominant system. This system was suggested by the Wikipedia page for Jacobi iteration.

```
1 def smooth(A, omega, nu, b, x0, tol=None):
2     """
3     smoothing function via -weighted Jacobi iteration
4     this is also a standard iterative method on a
5     diagonally dominant system
6
7     AA = np.array([[10., -1., 2., 0.],
8                   [-1., 11., -1., 3.],
9                   [2., -1., 10., -1.],
10                  [0.0, 3., -1., 8.]])
11
12     bb = np.array([6., 25., -11., 15.])
13     x00 = np.zeros_like(bb)
14     ans = smooth(AA, 1., 50, bb, x00)
15
16     converges after 24 iterations
17
18     if omega is changed to 2/3 in the above, converges in 35 iterations
19     """
20     if x0.ndim == 1:
21         x0 = np.expand_dims(x0, -1)
22     if b.ndim == 1:
23         b = np.expand_dims(b, -1)
24
25     x = x0.copy()
26     D = np.diag(A) # diagonal of system (as a Nx1 vector)
27     # must be same shape as b or will broadcast to a matrix under division
28     D = D.reshape(x.shape)
29
30     W = np.tril(A, k=-1) + np.triu(A, k=1) #deleted diagonal
31
32     for i in range(nu):
33         x = (1-omega)*x + ((omega*(b- (W@x))) / D)
34         if tol is not None and np.allclose(b, A@x, 1e-12):
35             break
36     else:
37         if tol is not None:
38             print("Warning, did not converge within tolerance", tol)
39
40     return x
```

## 2.3 MGCG (Full Method)

```
1 def mgcg(A, b, n_levels=0, interpolation_method=None, verbose=False):
2     """
3     Multigrid preconditioned Conjugate Gradient Method
4     This solves the system  $Ax=b$  for a particular system A.
5
6     if n_levels is 0 do unpreconditioned CG
7     """
8     M = int(np.log2(A.shape[0]))
9     # check that A is actually a power of 2 (no strategy otherwise)
10    N = A.shape[0]
11
12    assert 2**M == N
13
14    preconditioned = (n_levels != 0 and interpolation_method is not None)
15
16    if preconditioned:
17        # build *all* interpolation matrices and store
18        I = tuple((interpolation_method(M, n_levels, l)
19                    for l in range(n_levels, 1, -1)))
20        a = A
21        A_levels = list() # store all systems
22        for interp in I:
23            A_levels.append(a) # append last system matrix
24
25            # yer done if interp is 0x0 (only an issue if L is larger than M)
26            if not interp.size:
27                break
28
29            a = interp.T @ (a @ interp)
30
31        # now base_case
32        A_1 = a
33        A_levels = tuple(A_levels) # make static
34
35        # you may check np.allclose(Lchol, linalg.cholesky(A_1, lower=True))
36        Lchol, it_chol = cholesky(A_1)
37
38        Minv = lambda b: vcycle(n_levels, b, np.zeros_like(b), A_levels, I, Lchol)
39
40
41    else:
42        # if n_levels == 0 then do unpreconditioned CG
43        # i.e. preconditioner is the identity
44        Minv = lambda b: b
45
46    pcg_sol, its = pcg(A, b, Minv, return_iterations=True, verbose=verbose)
47
48    if preconditioned:
49        return pcg_sol, its, A_levels, I
50    else:
51        return pcg_sol, its, A, None
52
```

### 3 Results

When the program `mgcg.py` is run by itself, the following output is generated. Here, `interpolation_matrix` and `interpolation_matrix_2` refer to PCG in scenario #1 and #2, respectively, while `None` refers to unpreconditioned CG.

---

```
M= 11 L= 6
method: <function interpolation_matrix at 0x7f52b04bbc80>
  65 iterations
  time elapsed= 8.153933763504028
```

---

```
M= 11 L= 6
method: <function interpolation_matrix_2 at 0x7f52b04bbd08>
  41 iterations
  time elapsed= 4.9130401611328125
```

---

```
M= 11 L= 6
method: None
  512 iterations
  time elapsed= 3.9258689880371094
```

---

```
M= 12 L= 7
method: <function interpolation_matrix at 0x7f52b04bbc80>
  99 iterations
  time elapsed= 46.1121826171875
```

---

```
M= 12 L= 7
method: <function interpolation_matrix_2 at 0x7f52b04bbd08>
  61 iterations
  time elapsed= 28.609034299850464
```

---

```
M= 12 L= 7
method: None
  1024 iterations
  time elapsed= 29.122495651245117
```

---

```
M= 13 L= 8
method: <function interpolation_matrix at 0x7f52b04bbc80>
  158 iterations
  time elapsed= 258.28484988212585
```

---

```
M= 13 L= 8
method: <function interpolation_matrix_2 at 0x7f52b04bbd08>
  87 iterations
  time elapsed= 147.13387250900269
```

---

```
M= 13 L= 8
method: None
  2048 iterations
  time elapsed= 237.42299604415894
```

---

For clarity, these results are summarized in a table below.

### 3.1 Calculation of A-multiplies in CG/PCG

A fairly straightforward implementation of preconditioned conjugate gradient method is given below.

```
1  def pcg(A,b, Minv, tol=1e-8, x_init=None, return_iterations=False,
2      return_error=False, verbose=False):
3      """
4      preconditioned conjugate gradient method
5      solves Ax = b by preconditioning
6
7      INPUT:
8
9      A      - an NxN nd.array describing the system
10     b      - initial conditions (can be 1D-array or 2D row vector)
11     Minv    - preconditioner, which should be a function handle
12     tol     - stopping tolerance (returns sol if ||A*sol - b||_2 < tol )
13              (optional) default is 1e-8
14     x_init  - initial guess (default is None, in which case the zero vector
15              is used)
16
17     return_iterations (optional) return iteration count (default False)
18     return_error      (optional) return error (will be below tolerance
19                      if converged)
20
21     OUTPUT:
22     x          - solution (an Nx1 nd.array)
23     iterations  -(if return_iterations=True above) iterations to run
24     err         -(||A*sol-b|| of solution calculated error of residual
25     """
26
27     # make sure initial guess is a column vector ala matlab
28     if b.ndim == 1:
29         b = np.expand_dims(b,-1)
30
31     if x_init is None:
32         x = np.zeros_like(b) # default to zero vector as initial guess
33     else:
34         x = x_init
35
36     tol *= norm(b) # for stopping check (save some divisions)
37
38     r = b - A@x      # initial residual
39     z = Minv(r)       # residual of preconditioned system
40     p = z.copy()      # initial search direction
41     d = A@p           # initial A-projected search direction
42
43
44     for iterations in count(1):
45
46         alpha = np.vdot(r,z) / np.vdot(p,d)
47
48         x += alpha*p
```

```

49     r_new = r - alpha*d
50
51     # equivalent to norm(b - A@x) / norm(b)
52     err = norm(r_new)
53
54     if verbose:
55         print(iterations, err, sep='\t| ')
56
57     if err <= tol:
58         break
59
60     z_new = Minv(r_new)
61     beta = np.vdot(z_new,r_new) / np.vdot(r,z)
62
63     p = z_new + beta*p
64
65     d = A@p
66
67     r = r_new
68     z = z_new
69
70
71     #if err <= tol:
72     #    break
73
74     # return statement boogaloo
75     if return_iterations:
76
77         if return_error:
78             return x, iterations, err / norm(b)
79         else:
80             return x, iterations
81
82     elif return_error:
83         return x, err / norm(b)
84
85     else:
86         return x

```

Careful counting of the above shows there is exactly one  $A$ -multiplication and one  $M^{-1}$ -multiplication in initialization, and one of each during each iteration. By hypothesis, each  $M^{-1}$  (which is a v-cycle) corresponds to 5  $A$ -multiplies. Thus PCG requires  $1 + k + 5 * (1 + k) = \boxed{6+6k}$   $A$ -multiplies for  $k$  iterations. By this logic, we can see that CG will only require  $\boxed{1 + k}$   $A$ -multiplies, (since CG is just PCG where  $M^{-1}$  is identity).

### 3.2 Tables (Iteration Counts, Speedup Ratio, Work Ratio)

So solving our system (1) with unpreconditioned CG converges in the worst case, the exact size of the system. Either preconditioner is a massive improvement. Preconditioner #2 converges in fewer iterations than preconditioner #1.

Our tables suggest that implementation #2 is much more efficient than implementation #1, especially considering the number of  $A$ -multiplications per v-cycle can be drastically reduced (see section below).

	M = 11	M = 12	M = 13
CG	512	1024	2048
PCG #1	65	99	158
PCG #2	41	61	87

Table 1: Iteration counts of CG & PCG on our system

	M = 11	M = 12	M = 13
CG vs. PCG #1	7.877	10.343	12.962
CG vs. PCG #2	12.488	16.787	23.540

Table 2: Speedup ratios of CG & PCG on our system (approximate)

	M = 11	M = 12	M = 13
CG vs. PCG #1	1.295	1.708	2.148
CG vs. PCG #2	2.036	2.755	3.88

Table 3: Work ratios of CG & PCG on the system (approximate)



## 4 The Bonus Question Answered

### 4.1 The Answer

If MGCG is performed on the system  $A$  (1) using the interpolation method **SC2**, the optimal choice of level-depth  $L$  is **exactly 3** for any size  $M \geq 3$ . That is, the “standard” interpolation step in (SC1) is used to interpolate between  $A_L \rightarrow A_{L-1}$  and  $A_{L-2} \rightarrow A_{L-3} = A_1$ , and the “special” interpolation method unique to **SC2** is performed once in between. There are two reasons why, and they are clearly seen from a visual depiction of the multigrid systems of  $A$  using scenario #2.

Refer to **Figure 1**, which was generated with the code `bonus_demo.py`. Here we see the four finest systems in the multigrid when  $M = 7, L = 6$ , although this result is obviously true for any choice of  $M$ .

### 4.2 A loose justification

The result is clear: applying the modified interpolation step between levels  $A_{L-1}$  and  $A_{L-2}$  converts the system from its original structure (with an antidiagonal) to a tridiagonal matrix, and all further interpolations preserve tridiagonality. Recalling that our goal in the V-cycle is to restrict the size of the system until we arrive at a system that is reasonable to solve directly, it makes sense that we should immediately solve such a system. There will be no additional benefit from successive restriction of the system once it is tridiagonal.

bonus\_demo.py

```
1 #!/usr/bin/env python3
2
3 from mgcg import *
4 import matplotlib.pyplot as plt
5 import numpy as np
6
7 M = 7
8 levels = 6
9 A_L = make_system(M)
10 b = np.ones((2**M,1))/2**(M/2)
11
12
13 sol, its, As, Is = mgcg(A_L,b,n_levels=levels,
14                        interpolation_method=interpolation_matrix_2)
15
16 fig, ax = plt.subplots(2,2)
17
18 ax = ax.ravel()
19
20 for i, a in enumerate(ax):
21     ax[i].spy(As[i])
22     if i > 0:
23         ax[i].set_title(r'$A_{\mathscr{l}}$',\;\mathscr{l}=L - '$'+str(i),
24                        fontsize=20)
25     else:
26         ax[i].set_title(r'$A_L$', fontsize=20)
27     ax[i].axis('off')
28
29 fig.tight_layout()
30 plt.show()
```




Figure 1: A look at the four finest levels of multigrid under **SC2**. After the second interpolation step, all subsequent levels are tridiagonal (and thus easily solved by Cholesky factorization)

### 4.3 Empirical Support via Run times

Of course, since performing the downscaling of the system  $A$  isn't computationally "free", we can directly see the superiority of choosing *exactly* 3 levels. The following shows a session in ipython, where we perform the entire MGCG procedure on our system when  $M = 11$ , iterating on the number of levels  $L = 1, 2, \dots, 9$ . The average runtime over 3 runs is displayed for each  $L$  value.

```
In [1]: from mgcg import *
In [2]: M = 11
In [3]: A = make_system(M)
In [4]: b = make_initial_conditions(M)
In [5]: for L in range(1,10):
.....:     %timeit mgcg(A,b,L,interpolation_matrix_2)
.....:
1 loop, best of 3: 2.65 s per loop
1 loop, best of 3: 1.86 s per loop
1 loop, best of 3: 1.48 s per loop <-----
1 loop, best of 3: 2.79 s per loop
1 loop, best of 3: 3.96 s per loop
1 loop, best of 3: 4.99 s per loop
1 loop, best of 3: 6.84 s per loop
1 loop, best of 3: 7.58 s per loop
1 loop, best of 3: 7.9 s per loop
```

These times are depicted in the graph below:

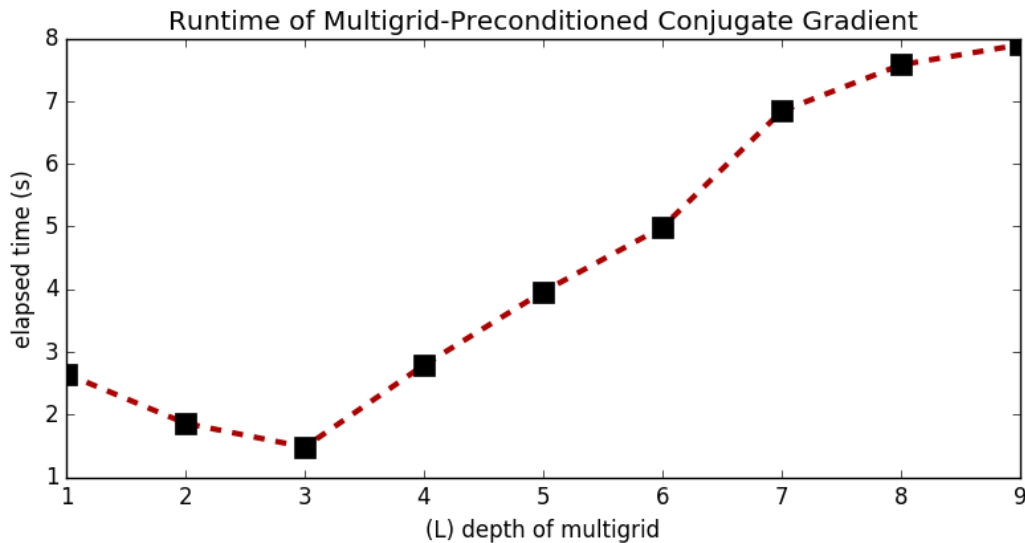


Figure 2: Measuring run time of solving  $Ax = b$  with MGCG when  $M = 11$ , and varying  $L$ . Ran on a laptop using unoptimized code.

## 4.4 An alternative justification

Figure 3 shows the effect of the preconditioners using both **SC1** and **SC2** on the initial conditions, as compared to the actual found solution of  $Ax = b$  with system (1). Note the progress made by **SC2** after 3 levels is both more than \*ever\* achieved by **SC1** when we compare how the initial conditions are transformed.

This picture also gives a strong (yet loose and possibly fallacious) rationalization for the particular nature of the novel step in **SC2**: After a single round of restriction, the L-1 level restriction/interpolation step *forces* the remaining solutions to be symmetric, with a crease in the middle (refer to implementation of **SC2** in the spec and also this report). It just so happens that the particular solution to our  $Ax = b$  is a very symmetrical solution, so (whatever the specifics), it is clear that the new step in **SC2** is exactly designed to exploit the *expected* symmetry of the solution to  $Ax = b$ , whereas (SC1) does not take advantage of the symmetry.




Figure 3: Measuring the proximity of a sample vector (the initial conditions and initial residual of PCG) to the eventual solution when  $M = 11$ .